

Smoothing Algorithm for n-gram Model

Author: *Zhiteng Li, Xiangge Huang, Huangji Wang*

Course: *Fall 2021, CS3602-2: Natural Language Processing*

Date: *October 12, 2021*

Contents

1	Program Basis	1
1.1	Background	1
1.2	Environment	1
1.3	Division	1
1.4	Model Description	2
1.5	Perplexity	2
2	N-gram Language Model	3
2.1	Foundation: 2-gram	3
2.2	Deepening: n-gram	3
3	Optimization: Smoothing	5
3.1	K-additive Smoothing	5
3.2	Jelinek-Mercer Smoothing (Interpolation)	5
3.3	Good Turing Discounting	7
3.4	Katz' s Back-off Model: Kenlm Model	8
3.5	Comparison	8
4	Discussion	9
4.1	Strengths	9
4.2	Weaknesses	9
5	Conclusion	10
A	Appendix for Codes	11
A.1	Kenlm Model in Ubuntu 18.08	11
A.2	Three Type of Smoothing methods	11

1. Program Basis

1.1. Background.

The language model construction is one of the most crucial parts in natural language processing, and now the most applied algorithm to solve this corresponding language term is to use the probability model to evaluate the natural language. According to the request for this project, we are expected to build our own n-gram language model on the basis of the word data and use relative smoothing way to optimize it.

Based on the analysis ahead, we use the corresponding method to achieve our 3-gram language model and evaluate the language part parameter perplexity.

1.2. Environment.

We use these environments and tools to finish this project.

- Python 3.8.7 and vscode in Windows 10
- Language Model: kenlm in Ubuntu 18.04
- Data Filter Tool: hash table

1.3. Division.

- Zhiteng Li:
 - Build n-gram model.
 - Design additive smooth, interpolation and Good-Turing discounting.
 - Code writing.
- Xiangge Huang:
 - Build 2-gram model.
 - Calculate the perplexity in the test set and train set in general way.
 - Paper writing.
- Huangji Wang:
 - Calculate the perplexity with this kenlm model.
 - Design additive smooth and test smooth interpolation.
 - Paper writing.

1.4. Model Description.

On the basis of the 2-gram model fundamental, we design one type of discounting algorithms and build our higher level n-gram language model and then use the train set to instantiate it. Then, after we get our n-gram and achieve the language word probability, we give the perplexity of the test set.

The basic n-gram language model can be described as:

$$P(w_k | w_{k-n+1}, \dots, w_{k-1})$$

This formula means when the current word probability is only decided by its former $n - 1$ words. The picture shows the 2-gram model as below.




Figure 1: 2-gram model

Obviously, in this 2-gram model, the latter words A, B and C emergence can be linked to and depend on the one word before it. It can be expressed by the conditional probability between these two words. So now we can have our own 2-gram model on this theory and deepen it into n-gram model to pursue more precision in expressing our language model.

1.5. Perplexity.

Additionally, after the building of the language probability model, the most important work is to evaluate the language data through our probability model, and we take perplexity to illustrate and quantify its demerits. Generally, the definition of perplexity is

$$PPL = \prod_{k=1}^K P(w_k | w_{k-n-1}^{k-1})^{-\frac{1}{K}}$$

This formula is based on the n-gram model, and also we can use kenlm library to get the perplexity easily. But because of the lacking of arpa form in our model, so the terminal outcome is calculated by our own perplexity calculation function.

2. N-gram Language Model

2.1. Foundation: 2-gram.

We build our language model from the easier 2-layer part, which is every word except the start single is only related to its former one word. So we can show the model and perplexity respectively.

$$P(w_k|w_{k-1}) \quad PPL = \prod_{k=1}^K P(w_k|w_{k-1})^{-\frac{1}{K}}$$

Use some words in the training data set as examples, the sentence

anarchism originated as a term of abuse

is the first vocabulary in this test set, and it will be transferred to these model

' < s > ' 'anarchism' 'originated' 'as' 'a' 'term' 'of' 'abuse'

so we can construct its corresponding probabilities

$$P(\text{anarchism} | < s >) \quad P(\text{originated} | \text{anarchism})$$

and others. Typically, the first part $P(< s >)$ is always thrown due to its value approaches 0 and it will lose the PPL precision.

Additionally, some probabilities used in 2-gram model are shown in Figure 2.

<s> anarchism	1.0
anarchism originated	0.004132231404958678
originated as	0.09417040358744394
as a	0.16765024236434656
a term	0.0013431549130990318
term of	0.02484152818228542
of abuse	0.00010881506670154329
abuse first	0.0019801980198019802
first used	0.006051437216338881
used against	0.0025797244634721992
against early	0.0001465845793022574
early working	0.00012238404112103782
working class	0.08112676056338028
class radicals	0.0007267441860465116
radicals including	0.01020408163265306
including the	0.2186314274499542
the diggers	7.024198363361781e-06
diggers of	0.047619047619047616

Figure 2: 2-gram probability

2.2. Deepening: n-gram.

After we build our 2-gram model, what we need to do is to deepen it and make it suitable for presenting our large data, so we build more deeper language model to solve this problem. And considering the data just consists of one long sentence, we believe the 3-gram model is enough to represent it, so we add one layer on the foundation of our 2-gram model. And Figure 3 exhibits some probabilities derived from this model.

<s> anarchism originated	1.0
anarchism originated as	1.0
originated as a	0.40476190476190477
as a term	0.0016281158769368964
a term of	0.12320916905444126
term of abuse	0.034482758620689655
of abuse first	0.019230769230769232
abuse first used	1.0
first used against	0.007142857142857143
used against early	0.02127659574468085
against early working	1.0
early working class	1.0
working class radicals	0.006944444444444444
class radicals including	0.5
radicals including the	1.0
including the diggers	0.0005984440454817474
the diggers of	0.16666666666666666
diggers of the	1.0

Figure 3: 3-gram probability

However, because of our language model is so large that the sparsity will occur easily, and according to the requirement of implementing smoothing algorithm, we transfer our class model expression to hash map and build the relative probabilities data. And the raw probabilities in 3-gram model can be

$$P(w_3|w_1w_2) = \frac{C(w_1w_2w_3)}{C(w_1w_2)}$$

And the smoothing algorithm mentioned below will solve the sparsity including unseen and low frequency. Also, considering our control group kenlm 3-gram language model, due to its probabilities log expression, we provide the equivalent form of *PPL* to make our instructions clear and scientific.

The original form and log equivalent form of *PPL* for 3-gram model is

$$PPL = \prod_{k=1}^K P(w_k|w_{k-1}w_{k-2})^{-\frac{1}{K}}$$

$$\log PPL = -\frac{1}{K} \sum_{k=1}^K \log P(w_k|w_{k-1}w_{k-2})$$

3. Optimization: Smoothing

The irrationality of data frequency is the fundamental reason why we adopt smoothing method, which has been exhibited in the raw probabilities ahead, so we use these ways to smooth our probability model.

Actually, we take three methods to realize the discounting algorithm. The first is k-additive smoothing. The second is Jelinek-Mercer smoothing. And the last is Good-Turing discounting. All the smoothing algorithms are based on the n-gram model.

3.1. K-additive Smoothing.

For normal additive smoothing, it sets $k = 1$. If $0 < k \leq 1$, the test result must be nicer. Thus, we try to traverse the test set and judge whether each 3-word phrase lies in the n-gram model. If the phrase is not in the model, we give it a low weight 0 and then it can be done by k-additive smoothing.

$$P(w_i|w_{i-n+1}^{i-1}) = \frac{k + C(w_{i-n+1}^{i-1})}{k|V| + \sum_{w_i} C(w_{i-n+1}^i)}, 0 < k \leq 1$$

Difficulties

Researchers have found that when the situation of $P(w_i) = 0$ occurs in k-additive smoothing, the result may be seriously bad. To deal with this problem, we firstly give each 3-word phrase one enough weight. However, if the 2-word phrase is not in the train set, our algorithm will leads to terrible error. Thus, we give each 2-word phrase that is not in the model 0 weight.

```
25     # pass in the initial value of para k
26     instance1 = Lidstone_smoothing(dev_data, k=0.5)
27
28     # smoothing
29     instance1.smoothing(hash2)
30
31     # calculate the PPL of test_set
32     res1 = instance1.calculate_PPL()
33     print(res1)
```

... 40.482017237459445

Figure 4: The PPL of K-additive smoothing.

3.2. Jelinek-Mercer Smoothing (Interpolation).

Actually, when we do k-additive smoothing, we rudely set two 3-word phrases that are not in the n-gram model the same weight and possibility. However, this result is obviously ambiguous. For example, "ZZZ" and "HAPPY" might not be totally equal in the same text. Thus, the Jelinek-Mercer Smoothing is to do interpolation in n-gram models. It can efficiently

handle with this problem.

$$P^*(w_i|w_{i-1}) = \lambda P(w_i|w_{i-1}) + (1 - \lambda)P(w_i)$$

In our 3-gram model, we hope to apply this method in our 3-gram model. Then, we apply the high order interpolation in our model. The key problem is to find good λ_n .

$$P^*(w_i|W_{i-n+1}^{i-1}) = \lambda_n P(w_i|W_{i-n+1}^{i-1}) + (1 - \lambda_n)P(w_i|W_{i-n+2}^{i-1})$$

To train λ_n , we use one good method called "Gradient Descent". Gradient descent is widely used in unconstrained optimization problems. For convex differential f ,

$$d_k \text{ is a descent direction} \Leftrightarrow d_k^T \nabla f(x_k) < 0$$

Algorithm 1 Descent Method

choose initial point $x_0 \in R^n$

repeat

 choose **descent direction** $d_k \in R^n$ and **step size** $t_k > 0$

$x_{k+1} = x_k + t_k d_k$ s.t. $f(x_{k+1}) < f(x_k)$

until stopping criterion is satisfied

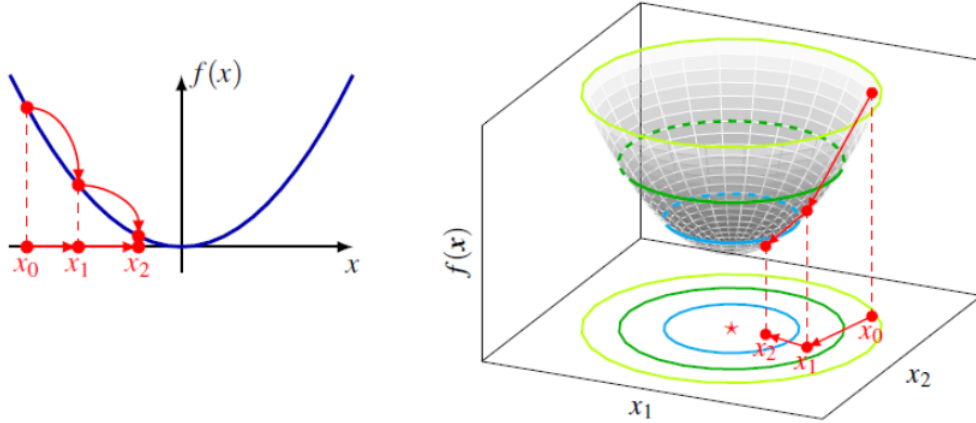


Figure 5: Principle of gradient descent

In our model, we calculate PPL by the formula:

$$PPL = (10^{-\sum \log(\lambda_2 P_3 + \lambda_1 (1 - \lambda_2) P_2 + (1 - \lambda_1) (1 - \lambda_2) P_1)})^{\frac{1}{\text{length}}}$$

Luckily, it's obvious that the function of $f(\lambda_1, \lambda_2)$ is convex because the Hessian matrix of $f(\lambda_1, \lambda_2)$ is positive definite ($0 \leq P_1, P_2, P_3 \leq 1$). It means we can use descent method directly and it will lead to perfect result.

Difficulties

Due to the reason that the function is discrete, if we want to build one complete function, the cost of gradient calculation must be terribly high. Thus, we apply **Expectation-maximization algorithm** in our gradient descent method to improve the efficiency. And

this idea leads to nice result. We firstly fix one parameter and do gradient descent in another parameter. Next, we fix this unfixed parameter and let another parameter find the gradient descent. We repeat these methods to find the local minimum parameters. Due to the reason that $f(\lambda_1, \lambda_2)$ is convex, the local minimum is actually the global minimum.




Figure 6: Example of EM algorithm

At last, the result is very nice due to the combination of gradient descent algorithm and EM algorithm.

```

1  # snd test the model on the test_set with paras lambda_1 and lambda_2
2
3  # pass in the initial values of lambda_1 and lambda_2
4  instance2 = Interpolation_smoothing(dev_data, test_data, lambda_1=0.5, lambda_2=0.5)
5
6  # train lambda_1 and lambda_2 on the dev_set
7  instance2.train_paras(hash1, hash2, hash3, dev_data, maxiter=1000, stepSize=0.05, tol=1e-5)
8
9  # calculate PPL on the test_set
10 res2 = instance2.calculate_PPL(hash1, hash2, hash3)
11 print("PPL of test_set: {}".format(res2))
12
...
-----
lambda_1: 0.5818039058034233
lambda_2: 0.18286565390627954
410.50279117052855

```

Figure 7: The optimal value of λ_1 and λ_2 and the PPL of our interpolation algorithm.

3.3. Good Turing Discounting.

The main idea is to redistribute the words of $r + 1$ times to the words of r times in the n -gram model. To finish this discounting algorithm, we should firstly modify N_r .

$$r^* = (r + 1) \frac{N_{r+1}}{N_r} \text{ and } d(r) = \frac{(r + 1)N_{r+1}}{rN_r} \text{ and } P(v : C(v) = r) = \frac{r^*}{N}$$

r	Symbols	N_r	rN_r	$P(x)$		r^*	Symbols	N_r	rN_r	$P(x)$
0	C, H	2	0	0		2	C, H	2	4	4/13
1	A, B, E, I	4	4	4/13		3/2	A, B, E, I	4	6	6/13
2	D, F, G	3	6	6/13	→	1	D, F, G	3	3	3/13
3	J	1	3	3/13		1	J	1	0	0
Total		13		1		Total		13		1

Figure 8: The example of Good Turing Discounting.

Based on these two previous algorithms, we finish this algorithm very quickly. The result is putted below.

```

40
47 # pass in the train_set, the test_set and the number of low frequency word
48 instance3 = discounting(hash2, test2, count=10)
49
50 # Good Turing Discounting
51 instance3.discounting()
52
53 # calculate the PPL of test_set
54 res3 = instance3.calculate_PPL()
55 print("ppl for discounting: {}".format(res3))
... ppl for discounting: 68.0130117421328

```

Figure 9: The PPL of Good Turing Discounting.

3.4. Katz' s Back-off Model: Kenlm Model.

Based on the existing model: kenlm model, we build the **arpa** file of the train set and put the test set into the model. And the result is clearly good.



```

1 import kenlm
2
3 model = kenlm.LanguageModel("test_train.bin")
4
5 f = open("test_set.txt", "r")
6
7 s = f.read()
8
9 print(model.perplexity(s))
10
11
12 os@ubuntu: ~/Desktop/kenlm/build$ python3 test.py
379.8103058534601

```

Figure 10: The PPL of kenlm model

3.5. Comparison.

Above all the result, we find that efficiency satisfies that (We don't consider K-additive smoothing because it isn't correct when it meets 0 possibility):

Good Turing Discounting > Kenlm Model > Interpolation

4. Discussion

The main goal of this lab is to understand n-gram model and language perplexity. By using different kinds of discounting algorithms, we get some efficient perplexities, and then choose the best of them according to the development data. For some details in this lab, we have queried class TA and got the corresponding response. Now, we will discuss the pros and cons of this model and explain the difficulties and solutions we have met and done.

4.1. Strengths.

1. Plenty Types of Smoothing and Nice Control Group

We have finished 3 types of smoothing algorithms and run them well. To test our improved results, we use the existing kenlm model to calculate the back-off PPL and compare it with our three PPL results. It shows that some of our PPL results are good while others are not good enough.

2. Use **dev-set** to Train Parameters

We make full use of three sets the teacher provided. The dev-set is well applied in our parameters training. λ_1 and λ_2 these two parameters have good convergence. Because the dev-set has no connection with the train-set, the parameters have excellent independence.

3. More Real Condition: 3-gram

The information of 3-gram is much larger than the information of 2-gram. It means the prefixes of 3-gram is more complex than 2-gram. The more information we get, the realer result we have. If we want the real condition, then n in n -gram model must lead to infinity, which is impossible in real situation. Thus, we use 3-gram that is easy to finish and calculate. The PPL is less than the PPL in 2-gram and 1-gram.

4. Good Function Gradation

We try to make all the codes simpler and cleaner to be understood. Thus, our functions in the codes has good hierarchy. Our codes are easy to debug and read.

4.2. Weaknesses.

1. Slow in Gradient Descent

The compute complexity of gradient descent is incredibly high. Thus when we execute the code, it runs very slow and often costs 6-8 minutes. Anyway, the result is precise.

2. Lack of 0 Possibility in K-additive Smoothing

We rudely set the items with 0 possibility. And in smoothing method we give it a k -weight. The result may be ambiguous due to the smoothing method. It means that the model is in bad conditions when it meets $P = 0\%$ item.

5. Conclusion

Briefly speaking, we have built our n-gram language model, implemented some smoothing algorithms and tested relative perplexities based on test data. The key point to this lab is to set a rational smoothing algorithm to let the language perplexities decrease so the whole model sparsity problem can be solved. So we use four traditional smoothing algorithm to finish this work, and the correlative outcomes have shown in part Smoothing. Additionally, we also have researched other smoothing algorithms like Absolutely discounting, while they are not presented in our major work.

In general, on the basis of the data provided, we construct our final 3-gram language model and then use smoothing way to decrease its low frequency. After that we add the train data to initialize our model and acquire its probabilities values. Also, through development data we choose the best smoothing algorithm in the combination. Finally, we reach the final goal returning the language perplexity on the test data and then analyze this whole project.

As the process of calculating the perplexities, we use kenlm library to aid and check the rationality of our results. Considering the library situation that it uses Backoff algorithm to build language model, we notice that our smoothing results are expected to be similar with it in numerical terms, and the before-smoothing results are supposed to be larger than it. And recalling our explanation of these smoothing perplexities, we can be sure that our smoothing algorithm is useful in the 3-gram language model.

A. Appendix for Codes

A.1. Kenlm Model in Ubuntu 18.08.

The code for build the arpa set:

```
bin/lmplz -o 3 -verbose header -text mytext/19-01-au.txt -arpa mylmmodel/test.arpa
```

```
1 import kenlm
model = kenlm.LanguageModel("test_train.bin")
3 f = open("test_set.txt", "r")
s = f.read()
5 print(model.perplexity(s))
```

Listing 1: Kenlm.py

A.2. Three Type of Smoothing methods.

We finish these methods in the **ipynb** format. The codes in **py** format are presented below.

```
1 # library
import math
3
# Lidstone Smoothing
5 class Lidstone_smoothing:
    def __init__(self, dev_data, k):
6
7         self.dev_data = dev_data
8
9
10        # log10 probability of each 3 words phase
11        self.hash = {}
12
13        # para, k=1 is Laplacian Smoothing
14        self.k = k
15
16        # len of dev_data
17        self.len = len(self.dev_data)
18
19    def smoothing(self, hash2):
20        for i in range(2, self.len):
21            # the previous two words
22            prefix = self.dev_data[i-2] + ' ' + self.dev_data[i-1]
23            # 3 words phase
24            tmp = self.dev_data[i-2] + ' ' + self.dev_data[i-1] + ' ' + self.dev_data[i]
25            # prefix not in hash2
26            if (prefix not in hash2):
27                pass
28            elif (tmp not in hash2[prefix]):
29                hash2[prefix][tmp] = 0
30
31        for key2 in hash2:
```

```

33         # number of str start in prefix "key2"
34         v = len(hash2[key2].keys())
35         m = sum(hash2[key2].values())
36         for key3 in hash2[key2]:
37             # Lidstone Smoothing
38             self.hash[key3] = math.log10((hash2[key2][key3] + self.k) / (m + v*self.k))
39
40     def calculate_PPL(self):
41         # calculate the PPL
42         ppl = 0
43
44         # number of 3 words phase
45         cnt = 0
46
47         for i in range(2, self.len):
48             # 3 words phase
49             tmp = self.dev_data[i-2] + ' ' + self.dev_data[i-1] + ' ' + self.dev_data[i]
50             if (tmp not in self.hash):
51                 continue
52             ppl += self.hash[tmp]
53             cnt += 1
54
55         ppl *= -1
56         ppl /= cnt
57         ppl = 10**ppl
58
59         return ppl
60
61     # Jelinek-Mercer Smoothing
62     class Interpolation_smoothing:
63         def __init__(self, dev_data, test_data, lambda_1, lambda_2):
64             self.dev_data = dev_data
65             self.test_data = test_data
66
67             # para lambda, learn from dev_set
68             self.lambda_1 = lambda_1
69             self.lambda_2 = lambda_2
70
71             # len of dev_data
72             self.len = len(self.dev_data)
73
74             # maxiter, stepSize and tol represent maxDepth of iteration, update step size and stop flag
75             # respectively
76             def train_paras(self, hash1, hash2, hash3, dev_data, maxiter, stepSize, tol):
77
78                 # gradient decent for lambda_2
79                 def gd2(lambda_1, lambda_2):
80                     res = 0
81
82                     # number of derivative
83                     cnt = 0
84
85                     for i in range(2, len(dev_data)):
86                         word3 = dev_data[i-2] + ' ' + dev_data[i-1] + ' ' + dev_data[i]
87                         word2 = dev_data[i-1] + ' ' + dev_data[i]
88                         word1 = dev_data[i]

```

```

87         # probability of different length phase
89         p3 = hash3[word3][1] if word3 in hash3 else 0
91         p2 = hash2[word2][1] if word2 in hash2 else 0
93         p1 = hash1[word1] if word1 in hash1 else 1
95
97         # sum the derivative of each phase
99         res += (p3-lambda_1*p2-(1-lambda_1)*p1)/(lambda_2*p3 + lambda_1*(1-lambda_2)*p2 + (1-
101             lambda_1)*(1-lambda_2)*p1)
103         cnt += 1
105
107         # get the avg of PPL derivative
109         res /= -cnt
111
113         return res
115
117     # gradient decent for lambda_1
119     def gd1(lambda_1, lambda_2):
121         res = 0
123
125         # number of derivative
127         cnt = 0
129
131         for i in range(2, len(dev_data)):
133             word3 = dev_data[i-2] + ' ' + dev_data[i-1] + ' ' + dev_data[i]
135             word2 = dev_data[i-1] + ' ' + dev_data[i]
137             word1 = dev_data[i]
139
141             # probability of different length phase
143             p3 = hash3[word3][1] if word3 in hash3 else 0
145             p2 = hash2[word2][1] if word2 in hash2 else 0
147             p1 = hash1[word1] if word1 in hash1 else 1
149
151             # sum the derivative of each phase
153             res += ((1-lambda_2)*p2-(1-lambda_2)*p1)/(lambda_2*p3 + lambda_1*(1-lambda_2)*p2 + (1-
155                 lambda_1)*(1-lambda_2)*p1)
157             cnt += 1
159
161             # get the avg of PPL derivative
163             res /= -cnt
165
167             return res
169
171     # fst train para lambda_1 and lambda2 on the dev_set
173
175     # EM algorithm - fst fix one var and update another, then fix another var and update the previous
177     one
179     while (gd2(self.lambda_1, self.lambda_2)**2 >= tol or gd1(self.lambda_1, self.lambda_2)**2 >= tol):
181         # update lambda_2
183         cnt = 0
185         gd = gd2(self.lambda_1, self.lambda_2)
187         while (cnt <= maxiter and gd**2 >= tol):
189             # update lambda_2 based on the gradient
191             self.lambda_2 -= gd*stepSize

```

```

141         gd = gd2(self.lambda_1, self.lambda_2)
        cnt += 1

143         # update lambda_1
        cnt = 0
145         gd = gd1(self.lambda_1, self.lambda_2)
        while (cnt <= maxiter and gd**2 >= tol):
147             # update lambda_1 based on the gradient
            self.lambda_1 -= gd*stepSize

149             gd = gd1(self.lambda_1, self.lambda_2)
            cnt += 1

151             gd = gd1(self.lambda_1, self.lambda_2)
            cnt += 1

153
155         print("-----")
        print("lambda_1: {}".format(self.lambda_1))
        print("lambda_2: {}".format(self.lambda_2))
157
159     def calculate_PPL(self, hash1, hash2, hash3):
        # calculate the PPL
161         ppl = 0

163         for i in range(2, self.len):
            word3 = self.test_data[i-2] + ' ' + self.test_data[i-1] + ' ' + self.test_data[i]
165             word2 = self.test_data[i-1] + ' ' + self.test_data[i]
            word1 = self.test_data[i]

167             # probability of different length phase
169             p3 = hash3[word3][1] if word3 in hash3 else 0
            p2 = hash2[word2][1] if word2 in hash2 else 0
171             p1 = hash1[word1][1] if word1 in hash1 else 1

173             # interpolation
            # lambda_1 and lambda_2 are learned from dev_set
175             ppl += math.log10(self.lambda_2*p3 + self.lambda_1*(1-self.lambda_2)*p2 + (1-self.lambda_1)*(1-
                self.lambda_2)*p1)

177         ppl *= -1
        ppl /= (self.len-2)
179         ppl = 10**ppl

181         return ppl

183 # Good Turing Discounting
class discounting:
185     def __init__(self, hash2, test2, count):
        # train_set
187         self.hash2 = hash2

189         # test_set
        self.test2 = test2

191         # define the number of low frequency word
193         self.numlow = count

```

```

195     # record probability of 3 word phase
    self.hash3 = {}
197
198 def discounting(self):
199     for key2 in hash2:
200         if (key2 not in self.test2):
201             continue
202
203     # find the maximum number of occurrences in train_set
    values = list(self.hash2[key2].values())
205     m = max(values)
206
207     # vocabulary size
    v = sum(self.hash2[key2].values())
209
210     # probability array
    p = [0 for _ in range(m+1)]
211
212     # record the number of words of the same occurrences
    cnt = [0 for _ in range(m+1)]
213
214     # record the number of words of the same occurrences
    for word3 in self.test2[key2]:
215         if (word3 not in self.hash2[key2]):
216             cnt[0] += 1
217         else:
218             cnt[self.hash2[key2][word3]] += 1
219
220     # Good Turing Discounting for lower frequency word
    n = min(m, self.numlow)
221
222     for i in range(n):
223         if (cnt[i] == 0):
224             p[i] = 0
225         else:
226             j = i+1
227             while (j < n-1 and cnt[j]==0):
228                 j += 1
229             p[i] = (j)*cnt[j]/cnt[i] / v
230
231     # the higher item retains the original probability
    if (n < m+1):
232         for i in range(n, m+1):
233             p[i] = cnt[i] / v
234
235     # normalize to ensure the sum of probability is 1
    ss = sum(p)
236
237     if (ss == 0):
238         continue
239     for i in range(m+1):
240         p[i] /= ss
241
242     # record probability of 3 word phase
    for word3 in self.test2[key2]:
243         if (word3 not in self.hash2[key2]):
244             self.hash3[word3] = p[0]
245         else:

```



```

251         self.hash3[word3] = p[self.hash2[key2][word3]]

253     def calculate_PPL(self):
254         # calculate the PPL
255         ppl = 0
256
257         cnt = 0
258
259         for key2 in self.test2:
260             for word3 in self.test2[key2]:
261                 if (word3 not in self.hash3 or self.hash3[word3]==0):
262                     continue
263                 ppl += math.log10(self.hash3[word3])
264                 cnt += 1
265
266         ppl /= -cnt
267         ppl = 10**ppl
268
269         return ppl
270
271 # align
272 length = 30
273 if __name__ == '__main__':
274
275     train_data = []
276
277     # f = open("hw1_dataset\\train_set.txt", "r")
278     with open(r"hw1_dataset\train_set.txt", "r") as f:
279         # read by line
280         for line in f:
281             train_data.append(line)
282
283     # split by whitespace
284     train_data = train_data[0].split(" ")
285
286     # begin and end symbol
287     train_data.insert(0, '<s>')
288     train_data.append('</s>')
289
290     dev_data = []
291
292     with open(r"hw1_dataset\dev_set.txt", 'r') as f:
293         # read by line
294         for line in f:
295             dev_data.append(line)
296
297     # split by whitespace
298     dev_data = dev_data[0].split(" ")
299
300     # insert begin and end symbols
301     dev_data.insert(0, '<s>')
302     dev_data.append('</s>')
303
304     test_data = []
305
306     with open(r"hw1_dataset\test_set.txt", 'r') as f:

```

```

307     # read by line
    for line in f:
309         test_data.append(line)

311 # split by whitespace
test_data = test_data[0].split(" ")
313
315 # insert begin and end symbols
test_data.insert(0, '<s>')
test_data.append('</s>')
317
319 # test the Lidstone Smoothing algorithm
319
321 # record 3 words phase with its prefix
hash2 = {}

323 # 2-gram
for i in range(1, len(train_data)):
325     # concatenate two words
    tmp = train_data[i-1] + ' ' + train_data[i]
327
    # record the previous word
329     if (tmp not in hash2):
        hash2[train_data[i-1] + ' ' + train_data[i]] = {}
331
333 # 3-gram
for i in range(2, len(train_data)):
    # the previous two words
335     prefix = train_data[i-2] + ' ' + train_data[i-1]
    tmp = train_data[i-2] + ' ' + train_data[i-1] + ' ' + train_data[i]
337     if (tmp not in hash2[prefix]):
        hash2[prefix][tmp] = 1
339     else:
        hash2[prefix][tmp] += 1
341
343 # pass in the initial value of para k
instance1 = Lidstone_smoothing(dev_data, k=0.5)

345 # smoothing
instance1.smoothing(hash2)
347
349 # calculate the PPL of test_set
res1 = instance1.calculate_PPL()
print(res1)
351

353 # fst train paras lambda_1 and lambda_2 on dev_set

355 # 1-gram
hash1 = {}
357
359 # 2-gram
hash2 = {}

361 # 3-gram
hash3 = {}

```

```

363
364 # 1-gram
365 for word in train_data:
366     if (word not in hash1):
367         hash1[word] = 1
368     else:
369         hash1[word] += 1
370
371 # 2-gram
372 for i in range(1, len(train_data)):
373     # concatenate two words
374     tmp = train_data[i-1] + ' ' + train_data[i]
375
376     # record the previous word
377     if (tmp not in hash2):
378         hash2[train_data[i-1] + ' ' + train_data[i]] = [train_data[i-1], 1]
379     else:
380         hash2[train_data[i-1] + ' ' + train_data[i]][1] += 1
381
382 # 3-gram
383 for i in range(2, len(train_data)):
384     # concatenate three words
385     tmp = train_data[i-2] + ' ' + train_data[i-1] + ' ' + train_data[i]
386
387     if (tmp not in hash3):
388         hash3[tmp] = [train_data[i-2] + ' ' + train_data[i-1], 1]
389     else:
390         hash3[tmp][1] += 1
391
392 # proportion
393 for key in hash3:
394     hash3[key][1] /= hash2[hash3[key][0]][1]
395
396 for key in hash2:
397     hash2[key][1] /= hash1[hash2[key][0]]
398
399 for key in hash1:
400     hash1[key] /= len(train_data)
401
402 # snd test the model on the test_set with paras lambda_1 and lambda_2
403
404 # pass in the initial values of lambda_1 and lambda_2
405 instance2 = Interpolation_smoothing(dev_data, test_data, lambda_1=0.5, lambda_2=0.5)
406
407 # train lambda_1 and lambda_2 on the dev_set
408 instance2.train_paras(hash1, hash2, hash3, dev_data, maxiter=1000, stepSize=0.05, tol=1e-5)
409
410 # calculate PPL on the test_set
411 res2 = instance2.calculate_PPL(hash1, hash2, hash3)
412 print("PPL of test_set: {}".format(res2))
413
414 # test the Good Turing Discounting algorithm
415
416 # record 3 words phase with its prefix
417 hash2 = {}

```

```

419 # 2-gram
421 for i in range( len(train_data)-1):
423     # concatenate two words
423     tmp = train_data[i]
425
425     # record the previous word
425     if (tmp not in hash2):
427         hash2[train_data[i]] = {}
429
429 # 3-gram
429 for i in range(1, len(train_data)):
431     # the previous two words
431     prefix = train_data[i-1]
433     tmp = train_data[i-1] + ' ' + train_data[i]
433     if (tmp not in hash2[prefix]):
435         hash2[prefix][tmp] = 1
435     else:
437         hash2[prefix][tmp] += 1
439
439 # record 3 words phase with its prefix
439 test2 = {}
441
441 # 2-gram
443 for i in range( len(test_data)-1):
445     # concatenate two words
445     tmp = test_data[i]
447
447     # record the previous word
447     if (tmp not in test2):
449         test2[test_data[i]] = {}
449
451 # 3-gram
451 for i in range(1, len(test_data)):
453     # the previous two words
453     prefix = test_data[i-1]
455     tmp = test_data[i-1] + ' ' + test_data[i]
455     if (tmp not in test2[prefix]):
457         test2[prefix][tmp] = 1
457     else:
459         test2[prefix][tmp] += 1
459
461 # pass in the train_set, the test_set and the number of low frequency word
461 instance3 = discounting(hash2, test2, count=10)
463
463 # Good Turing Discounting
465 instance3.discounting()
467
467 # calculate the PPL of test_set
467 res3 = instance3.calculate_PPL()
469 print("ppl for discounting: {}".format(res3))

```

Listing 2: n-gram-model.py